

```
In [1]:
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels
import statsmodels.api as sm
import sklearn
from sklearn.model_selection import train_test_split
```

```
In [2]:
adv=pd.read_csv("https://raw.githubusercontent.com/rkmishracs/dataset/main/advertising.csv")
```

```
In [3]:
adv.head()
```

Out[3]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

```
In [4]:
adv.shape
```

Out[4]:

(200, 4)

```
In [5]:
adv.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0    TV          200 non-null   float64
 1   Radio       200 non-null   float64
 2  Newspaper   200 non-null   float64
 3   Sales       200 non-null   float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
In [6]:
adv.describe()
```

Out[6]:


	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

In [7]:

```
sns.regplot(x='TV', y='Sales', data=adv)
```

Out[7]:

<AxesSubplot:xlabel='TV', ylabel='Sales'>




In [8]:

```
sns.regplot(x='Radio', y='Sales', data=adv)
```

Out[8]:

<AxesSubplot:xlabel='Radio', ylabel='Sales'>




In [9]:

```
sns.regplot(x='Newspaper', y='Sales', data=adv)
```

Out[9]:

<AxesSubplot:xlabel='Newspaper', ylabel='Sales'>




In [10]:

```
sns.pairplot(data=adv, x_vars=['TV', 'Radio', 'Newspaper'], y_vars='Sales')
```

Out[10]:

<seaborn.axisgrid.PairGrid at 0x7f8088206b80>



In [11]:

```
adv.corr()
```

Out[11]:


	TV	Radio	Newspaper	Sales
TV	1.000000	0.054809	0.056648	0.901208
Radio	0.054809	1.000000	0.354104	0.349631
Newspaper	0.056648	0.354104	1.000000	0.157960
Sales	0.901208	0.349631	0.157960	1.000000

In [12]:

```
sns.heatmap(adv.corr())
```

Out[12]:

<AxesSubplot:>




In [13]:

```
sns.heatmap(adv.corr(), annot=True)
```

Out[13]:

<AxesSubplot:>



In [14]:

```
X=adv[ 'TV' ]  
y=adv[ 'Sales' ]
```

In [15]:

```
X_train, X_test, y_train, y_test= train_test_split(X,y, train_size=0.70, random_state=100)
```

In [16]:

```
X_train_sm=sm.add_constant(X_train)  
X_train_sm.head()
```

/Users/admin/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/tsatools.py:142: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only
x = pd.concat(x[:,order], 1)

Out[16]:

	const	TV
74	1.0	213.4
3	1.0	151.5
185	1.0	205.0
26	1.0	142.9
90	1.0	134.3

In [18]:

```
lr=sm.OLS(y_train, X_train_sm)
lr_model=lr.fit()
lr_model.params
```

Out[18]:

```
const    6.948683
TV       0.054546
dtype: float64
```

In [19]:

```
lr_model.summary()
```

Out[19]:

OLS Regression Results

Dep. Variable:	Sales	R-squared:	0.816
Model:	OLS	Adj. R-squared:	0.814
Method:	Least Squares	F-statistic:	611.2
Date:	Wed, 25 Jan 2023	Prob (F-statistic):	1.52e-52
Time:	14:11:22	Log-Likelihood:	-321.12
No. Observations:	140	AIC:	646.2
Df Residuals:	138	BIC:	652.1
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	6.9487	0.385	18.068	0.000	6.188	7.709
TV	0.0545	0.002	24.722	0.000	0.050	0.059

Omnibus:	0.027	Durbin-Watson:	2.196
Prob(Omnibus):	0.987	Jarque-Bera (JB):	0.150
Skew:	-0.006	Prob(JB):	0.928
Kurtosis:	2.840	Cond. No.	328.

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [20]:

```
plt.scatter(X_train, y_train)
```


Out[20]:

<matplotlib.collections.PathCollection at 0x7f80a33aebb0>



In [29]:

```
plt.scatter(X_train, y_train)
plt.plot(X_train, 6.948+0.054*X_train, 'r')
plt.show()
```



In [30]:

```
y_train_pred=lr_model.predict(X_train_sm)
y_train_pred
```

Out[30]:

```
74      18.588747
3       15.212365
185     18.130563
26      14.743271
90      14.274178
...
87      12.986898
103     17.197830
67      14.546907
24      10.346884
8        7.417777
Length: 140, dtype: float64
```

In [23]:

```
residual=y_train-y_train_pred
residual
```

Out[23]:

```
74      -1.588747
3        1.287635
185     4.469437
26        0.256729
90     -0.274178
...
87        3.013102
103     2.502170
67     -1.146907
24     -0.646884
8      -2.617777
Length: 140, dtype: float64
```

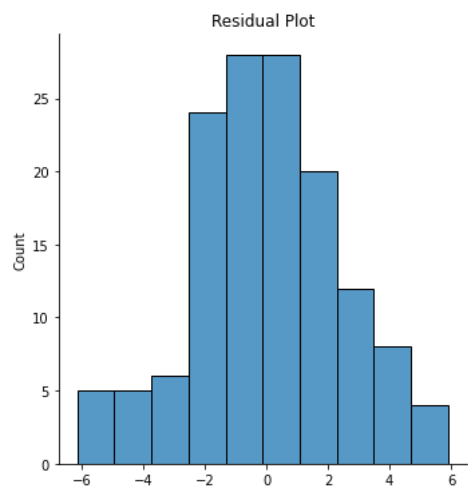
In [27]:

```
plt.figure()  
sns.displot(residual)  
plt.title("Residual Plot")
```

Out[27]:


Text(0.5, 1.0, 'Residual Plot')

<Figure size 432x288 with 0 Axes>



In [26]:

```
plt.scatter(X_train, residual)  
plt.show()
```



In []: